


## Old-value vs. correct answers (controlled dialogues, two models)


### (a) Still decodable

current-value probe score




### (b) Final layer

attention share



### (c) Attention drifts late

failures minus correct, length-adj.



## With vs. without conflict

### (d) Drift across models

six models, 95% intervals

