



A conflicting update pulls the read toward the old value in every model

What is compared · the same dialogue with vs. without a conflicting update

without a conflict




with a conflict




a Attention

shift toward the old value →



b Query-key match

shift toward the old value →



Same direction in all six models (ticks: 95% intervals).

c When the model actually fails

extra attention to old values on failures →



Shared pull, different failure signature:
the bias reaches the answer in one family,
not the other.